

Normalized 8-point Algorithm (DLT at the end)

- Given 8 or more corresponding points, we can estimate the fundamental matrix using the 8-point.
- The DLT algorithm involves solving a linear system of equations to estimate the fundamental matrix.

$$\begin{bmatrix} x'_m & y'_m & 1 \end{bmatrix} \begin{bmatrix} f_1 & f_2 & f_3 \\ f_4 & f_5 & f_6 \\ f_7 & f_8 & f_9 \end{bmatrix} \begin{bmatrix} x_m \\ y_m \\ 1 \end{bmatrix} = 0$$

$$\begin{bmatrix} x_1 x'_1 & x_1 y'_1 & x_1 & y_1 x'_1 & y_1 y'_1 & y_1 & x'_1 & y'_1 & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_m x'_m & x_m y'_m & x_m & y_m x'_m & y_m y'_m & y_m & x'_m & y'_m & 1 \end{bmatrix} \begin{bmatrix} f_1 \\ f_2 \\ f_3 \\ f_4 \\ f_5 \\ f_6 \\ f_7 \\ f_8 \\ f_9 \end{bmatrix} = 0$$

$$\mathbf{A}\mathbf{X} = \mathbf{0}$$

Algorithm Steps

- Normalize the points (since we are working in image coordinates)
 - Let the matrix used for normalizing the points in both cameras be \mathbf{T} and \mathbf{T}'
- Construct the matrix \mathbf{A} . Find SVD of \mathbf{A} , and using that, find approximate value $\hat{\mathbf{F}}$
- To find \mathbf{F} , enforce the rank-2 constraint on $\hat{\mathbf{F}}$ using SVD

$$\min_F \|F - \hat{F}\|_F, \text{ subject to } \det F = 0$$

$$F = U \begin{bmatrix} \Sigma_1 & 0 & 0 \\ 0 & \Sigma_2 & 0 \\ 0 & 0 & 0 \end{bmatrix} V^\top$$

- Denormalize the fundamental matrix to get the final fundamental matrix $F_{\text{final}} = (\mathbf{T}')^\top F \mathbf{T}$
- For DLT, every point generates 2 pairs $[x_1 \ y_1 \ 1 \ 0 \ 0 \ 0 \ -x_1 x'_1 \ -x'_1 y'_1 \ -x'_1]$, $[0 \ 0 \ 0 \ x_1 \ y_1 \ 1 \ -y_1 x'_1 \ -y'_1 y'_1 \ -y'_1]$ in \mathbf{A} , so you only need 4 corresp. points. Baki same algo only (except for rank 2 constraint)

Image Processing (just do convolution ez)

- Filtering: Form a new image whose pixel values are a combination of the original pixel values, to get useful info out of them.
- **Linear Filter:** Replace each pixel by a linear combination (a weighted sum) of its neighbors, the linear combination is the kernel.
- Cross Correlation:
 - Let F be the image, H be the kernel (of size $2k + 1 \times 2k + 1$), and G be the output image
 - $G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$ ($G = F \otimes H$)
 - Can think of it as a “dot product” between local neighborhood and kernel for each pixel
- Convolution:
 - Kernel is flipped (horizontally and vertically) before applying cross-correlation
 - $G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$ ($G = F * H$)
 - It is commutative, and associative.
- Averaging/Blur Filter: $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$, Sharpening Filter: $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$
- **Gaussian Filter:** Used for blurring, noise reduction, etc. Removes high frequency components from the image (low-pass filter)
 - Convolution with self is another Gaussian
 - Convolving twice with Gaussian kernel of width σ = convolving once with kernel of width $\sigma\sqrt{2}$
- Original - Blurred = High-pass filter, Original + α High-pass = Sharpened

Edge Detection

Convert a 2D image into a set of curves - Extracts salient features of the scene, More compact than pixels. - Edge is a place of rapid change in the image intensity function

How to take derivative for digital image? take discrete derivative $\frac{\partial f}{\partial x}[x, y] \approx F[x+1, y] - F[x, y]$. Gradient of an image is ∇f . Edge strength is $\|\nabla f\|$ and gradient direction is θ

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right], \|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}, \theta = \arctan\left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x}\right)$$

Noise in the image might cause too many rapid changes, so we smooth the image before taking the derivative. $\frac{d}{dx}(f * h) = f * \frac{d}{dx}h$. For 1D Gaussian,

$$G_\sigma(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}, G'_\sigma(x) = -\frac{1}{\sigma} \left(\frac{x}{\sigma}\right) G_\sigma(x)$$

Harris corner detection for feature detection

- Key Idea: Small shifts in the window should cause large changes in the intensity
- Compute image gradients for each pixel over a small region W . Let the image gradients be I_x, I_y .
- Subtract mean from each image gradient to remove the constant intensity
- Compute the covariance

$$H = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix}$$

This comes from second order Taylor expansion of $E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$, where x, y are the points in the window, and $w(x, y)$ is the window function.

- Compute eigenvectors and eigenvalues $Hx = \lambda x$
 - x_{\max} is the direction of the largest change in intensity
 - x_{\min} is the direction of the smallest change in intensity
 - $\lambda_{\max}, \lambda_{\min}$ are the amount of increase in direction of x_{\max} and x_{\min} respectively
- Depending on the value of λ_{\max} and λ_{\min} , we can classify the point as a corner, edge, or flat region
 - λ_{\max} is large, λ_{\min} is small \rightarrow Corners. λ_{\max} and λ_{\min} are large \rightarrow Edge. λ_{\max} and λ_{\min} are small \rightarrow Flat region
- Choose points where λ_{\min} is local maxima. Normally, we just approximate λ_{\min} with Harris operator, which is usually much easier to compute. one common choice is $\frac{\det(H)}{\text{trace}(H)} > \text{threshold}$

Feature Description

- Come up with a descriptor for each point, find similar descriptors in another image, and match them
- Invariance: Descriptor shouldn't change even if image is transformed
- Discriminability: Descriptor should be highly unique for each point
- Harris Detector is not invariant to scale, orientation, and illumination. SIFT is invariant to all of these.
- Automatic scale selection: Find scale that gives local maxima of some function f in both position and scale.
- From edges to blobs: Convolve image with Gaussian of different scales, and find local maxima in scale space

SIFT (Scale Invariant Feature Transform) for feature description

- Take 16x16 square window around detected feature
- Compute edge orientation (angle of the gradient - 90°) for each pixel
- Throw out weak edges (threshold gradient magnitude)
- Create histogram of surviving edge orientations
- Can handle changes in viewpoint (up to about 60 degree out of plane rotation)
- Can handle significant changes in illumination (sometimes even day vs. night)
- Pretty fast—hard to make real-time, but can run in <1s for moderate image sizes

RANSAC

- Randomly sample a subset of points, fit a model to that subset (Take 4 points, compute homography using DLT)
- Find the number of inliers for that model $[d(p'i, Hpi) < \epsilon]$
- Repeat for a large number of iterations, and choose the model with the most inliers (Re-compute least-squares H estimate using all of the inliers)
- Robustness : If there are G proportion of inliers, and our model need P pairs, Probability of picking P inliers = G^P , Probability of not picking set of inliers after N iterations = $(1 - G^P)^N$
- If there are multiple structures, RANSAC won't out of the box due to unknown number of structures, inlier set for each structure, and noise affecting the structures.